



## Classes

# Auto Steer

Auto Steering allows you to automatically generate feature interventions based on natural language descriptions of desired behaviors. This provides an easy way to steer model outputs without manually selecting features.

## Basic Usage

The simplest way to use Auto Steering is with the `AutoSteer` method:

Code

Output

```
# Create automatic feature edits for desired behavior
edits = client.features.AutoSteer(
    specification="be funny", # Natural language description
    model=variant, # Model variant to use
)

# Apply the edits to your variant
variant.set(edits)

# The model will now attempt to be funnier
response = client.chat.completions.create(
    messages=[{"role": "user", "content": "Tell me a story"}],
    model=variant
)
print(response.choices[0].message["content"])
```

## How It Works

Under the hood, Auto Steer:



1. Generates contrastive examples of content with and without the desired behavior
2. [Classes](#) > [Auto Steer](#)  
Identifies the most relevant features that distinguish the behavior
3. Determines optimal feature values to encourage the desired behavior
4. Creates a set of feature edits that can be applied to the model

## Advanced Usage

### Combining with Manual Edits

Auto Steer edits can be combined with manual feature interventions:

```
# Generate automatic edits
auto_edits = client.features.AutoSteer(
    specification="be professional",
    model=variant
)

# Combine with manual feature edits
variant.set({
    **auto_edits,
    manual_feature: 0.8
})
```

### Using with Conditionals

Auto Steer can be used within conditional statements:

```
# Generate edits for desired behavior
funny_edits = client.features.AutoSteer(
    specification="be funny",
    model=variant
)

# Apply edits conditionally
variant.set_when(context_feature > 0.5, funny_edits)
```

# Best Practices



Use clear, specific behavior descriptions

[Classes](#) > [Auto Steer](#)  
Test generated edits to ensure desired results

Consider combining multiple Auto Steer edits for complex behaviors

Adjust number of features based on steering precision needed

Use with conditionals for context-aware behavior

## API Reference

### AutoSteer

Generate automatic feature edits based on natural language description.

#### Parameters:

**specification** `str` *required*

Natural language description of desired behavior

**model** `Union[str, Variant]` *required*

Model or variant to use for generating edits

#### Returns:

**edits** `dict[Feature, float]`

Dictionary mapping features to their target values

#### Example:

##### Auto Steer

```
edits = client.features.AutoSteer(  
    specification="be more creative",
```



model=variant

Classes > Auto Steer

< Conditionals

Powered by Mintlify



## Classes

# Chat Completions

Documentation for the Goodfire Chat API

The Chat API provides methods for interacting with Goodfire's language models in a chat format. The base chat interface is OpenAI-compatible. It supports both streaming and non-streaming completions, as well as logits computation.

Once you have a [model variant](#), you can use it to create chat completions.

## Examples

### Basic Chat Completion

Code

Output

```
# Initialize the client
client = goodfire.Client(
    '{YOUR_API_KEY}',
)

# Create a non-streaming completion
response = client.chat.completions.create(
    messages=[
        {"role": "user", "content": "What is the capital of France?"}
    ],
    model="meta-llama/Llama-3.3-70B-Instruct"
)

print(response.choices[0].message["content"])
```

# Streaming Chat Completion



Streaming Chat Completion

Output Poem

Classes > Chat Completions

# Stream the response tokens

```
for chunk in client.chat.completions.create(
    messages=[
        {"role": "user", "content": "Write a short poem"}
    ],
    model="meta-llama/Llama-3.3-70B-Instruct",
    stream=True
):
    print(chunk.choices[0].delta.content, end="")
```

## Using with Model Variants

Using with Model Variants

Pirately Output

```
# Create a variant with feature modifications
variant = goodfire.Variant("meta-llama/Llama-3.3-70B-Instruct")
pirate_features = client.features.search(
    "talk like a pirate",
    model=variant,
    top_k=1
)
variant.set(pirate_features[0], 0.5)

# Use the variant in chat completion
for token in client.chat.completions.create(
    messages=[
        {"role": "user", "content": "Tell me about the ocean."}
    ],
    model=variant,
    stream=True,
    max_completion_tokens=100,
):
    print(token.choices[0].delta.content, end="")
```

## Computing Token Probabilities



```
logits = client.chat.logits(  
    messages=[  
        {"role": "user", "content": "The capital of France is"}  
    ],  
    model="meta-llama/Llama-3.3-70B-Instruct",  
    filter_vocabulary=["Paris", "London", "Berlin"]  
)  
print(logits.logits)
```

## Methods

### create()

Create a chat completion with the model.

#### Parameters:

**messages** `list[ChatMessage]` **required**

List of messages in the conversation. Each message should have `role` ("user", "assistant", or "system") and `content` fields.

**model** `Union[str, VariantInterface]` **required**

Model identifier or variant to use for completion

**stream** `bool` `default: "False"`


Whether to stream the response tokens

**max\_completion\_tokens** `Optional[int]` `default: "2048"`

Maximum number of tokens to generate

**top\_p** `float` `default: "0.9"`

Nucleus sampling parameter

 **temperature** float default: "0.6"

Sampling temperature

---

[Classes](#) > **Chat Completions**

---

**stop** Optional[Union[str, list[str]]]

Sequences where the API will stop generating further tokens

**seed** Optional[int] default: "42"

Random seed for reproducible outputs

**system\_prompt** str

System prompt to prepend to the conversation

---

Returns:

If `stream=False`: `ChatCompletion` object

If `stream=True`: Iterator of `StreamingChatCompletionChunk` objects

Examples:

Non-streaming completion:

#### Basic Chat Completion

```
response = client.chat.completions.create(
    messages=[
        {"role": "user", "content": "What is the capital of France?"}
    ],
    model="meta-llama/Llama-3.3-70B-Instruct"
)
print(response.choices[0].message["content"])
```

Streaming completion:

#### Streaming Chat Completion



```

for chunk in client.chat.completions.create(
    messages=[
        {"role": "user", "content": "Write a short poem"}
    ],
    model="meta-llama/Llama-3.3-70B-Instruct",
    stream=True
):
    print(chunk.choices[0].delta.content, end="")

```

## logits()

Compute token probabilities for the next token in a conversation.

### Parameters:

**messages** List[ChatMessage] required

List of messages in the conversation

**model** Union[str, VariantInterface] required

Model identifier or variant to use

**top\_k** Optional[int]

Limit response to top K most likely tokens


**filter\_vocabulary** Optional[list[str]]

List of tokens to compute probabilities for

**Returns:** `LogitsResponse` containing token probabilities

### Example:

#### Token Probabilities



```
logits = client.chat.logits(  
    messages=[  
        {"role": "user", "content": "The capital of France is"}  
    ],  
    model="meta-llama/Llama-3.3-70B-Instruct",  
    filter_vocabulary=["Paris", "London", "Berlin"]  
)  
print(logits.logits)
```

## Response Objects

### ChatCompletion

Response from a non-streaming chat completion.

#### Properties:

**id** str

Unique identifier for the completion

---

**object** str

Object type identifier

---

**created** Optional[int]

Unix timestamp of when the completion was created

---

**model** str

ID of the model used

---

**system\_fingerprint** str

System fingerprint for the completion

---

**choices** list[ChatCompletionChoice]



## StreamingChatCompletionChunk

Individual chunk from a streaming chat completion.

### Properties:

**id** str

Unique identifier for the completion

**object** str

Object type identifier

**created** Optional[int]

Unix timestamp of when the chunk was created

**model** str

ID of the model used

**system\_fingerprint** str

System fingerprint for the completion


**choices** list[StreamingChoice]

List of completion choices in this chunk

## LogitsResponse

Response from a logits computation request.

### Properties:

 `dict[str, float]`

Dictionary mapping tokens to their probabilities

---

[Classes](#) > **Chat Completions**

---

[< Client Reference](#)

**Variants** >

---

Powered by Mintlify



## Classes

# Client Reference

## Client

The Goodfire SDK provides both synchronous and asynchronous clients for interacting with the Goodfire API.

The examples provided in this guide are for the synchronous client. The asynchronous client is identical but with `await` added to the function calls.

We've broken down the SDK into a few key concepts:

[Chat Completions](#) - For interacting with language models

[Variants](#) - For managing model edits & variants

[Features](#) - For working with interpretable features

[Conditionals](#) - For working with conditional feature interventions

[AutoSteering](#) - For steering models with natural language

## Initialization

Sync

Async

```
from goodfire import Client

client = Client("your-api-key")
```



Powered by Mintlify

---

Classes > Client Reference

---



## Classes

# Conditionals

Documentation for working with conditional feature interventions

Conditionals allow you to define dynamic feature interventions that are applied based on the activation patterns of other features during model inference. This enables creating more sophisticated steering behaviors that respond to the content being generated.

Before using the Conditionals API, you'll need to find the features you want to intervene on, and a model variant

## Examples

### Basic Conditional Intervention

Apply pirate-themed features only when whale-related content is detected:

#### Basic Conditional Intervention

#### Output text

```
variant.reset()
# Find relevant features
whale_feature = client.features.search(
    "whales", model=variant, top_k=1
)

pirate_features = client.features.search(
    "talk like a pirate", model=variant, top_k=5
)

# Set up conditional intervention
variant.set_when(whale_feature > 0.75, {
    pirate_features[0]: 0.4
})
```



```
# The model will now talk like a pirate when discussing whales
response = client.chat.completions.create(
    messages=[{"role": "user", "content": "Tell me about whales."}],
    model=variant
)

print(response.choices[0].message["content"])
```

## Aborting Generation

Stop generation if certain content is detected:

### Abort example

### Abort output

```
# Abort if whale features are too strong
variant.abort_when(whale_feature > 0.75)

try:
    response = client.chat.completions.create(
        messages=[{"role": "user", "content": "Tell me about whales."}],
        model=variant
    )
except goodfire.exceptions.InferenceAbortedException:
    print("Generation aborted due to whale content")
```

## Auto-Generated Conditionals

Use natural language to automatically generate conditional statements:

### Auto-Generated Conditionals example

```
#create a variant
variant = goodfire.Variant("meta-llama/Llama-3.3-70B-Instruct")
# Generate conditional based on description - this will create conditions for k
conditional= client.features.AutoConditional(
    "when the model talks about whales and penguins",
    model=variant
)
```



```

# Get pirate feature
pirate_feature = client.features.search(
    "talk like a pirate", model=variant, top_k=1
)
Classes > Conditionals
# Make the model talk like a pirate when it talks about both whales and penguins
variant.set_when(conditional, {
    pirate_feature[0]: 0.9
})

response = client.chat.completions.create(
    messages=[{"role": "user", "content": "Tell me about whales and penguins!"}]
    model=variant
)

print(response.choices[0].message["content"])

```

## Creating Conditionals

### Comparison Operators

You can create conditionals by comparing features or feature groups with numeric values or other features using standard comparison operators. This creates a Conditional object that can be used in steering behaviors.

```

# Compare feature to numeric value
condition = feature > 0.75

# Compare feature group to numeric value
condition = feature_group >= 0.5

# Compare features to each other
condition = feature1 < feature2

```

Supported operators:

`==` (equal)

`!=` (not equal)

`<` (less than)



`<=` (less than or equal)

`>` (greater than)

`<=` (greater than or equal)

Classes > Conditionals

## Logical Operators

Multiple conditions can be combined using logical operators to create a

### ConditionalGroup:

```
# AND operator
condition = (feature1 > 0.5) & (feature2 < 0.3)

# OR operator
condition = (feature1 > 0.5) | (feature2 > 0.5)
```

## Using Conditionals

### set\_when()

Apply feature interventions when a condition is met.

#### Parameters:

**condition** `ConditionalGroup` **required**

The ConditionalGroup that triggers the intervention

**values** `Union[FeatureEdits, dict[Union[Feature, FeatureGroup], float]]` **required**

Feature edits to apply when condition is met

Returns: None

#### Example:

```
# Set pirate features when whale features are detected
variant.set_when(whale_feature > 0.75, {
```



```
pirate_features[0]: 0.5
```

Classes > Conditionals

## abort\_when()

Abort inference when a condition is met by raising an `InferenceAbortedException`.

Parameters:

**condition** `ConditionalGroup` **required**

The `ConditionalGroup` that triggers the abort

---

Returns: None

Example:

```
# Abort if whale features are too strong
variant.abort_when(whale_feature > 0.75)

try:
    response = client.chat.completions.create(
        messages=[{"role": "user", "content": "Tell me about whales."}],
        model=variant
    )
except goodfire.exceptions.InferenceAbortedException:
    print("Generation aborted due to whale content")
```

## handle\_when()

Register a custom handler function to be called when a condition is met.

Parameters:

**condition** `ConditionalGroup` **required**

The `ConditionalGroup` that triggers the handler

---

 **handler** Callable[[InferenceContext], None] required

Function that takes an InferenceContext and returns None

Classes > **Conditionals**

---

Returns: None

Example:

```
def custom_handler(context: InferenceContext):  
    # Custom handling logic  
    pass  
  
variant.handle_when(whale_feature > 0.5, custom_handler)
```

## AutoConditional

The AutoConditional utility helps automatically generate conditional statements based on natural language descriptions.

Parameters:

**specification** str required

Natural language description of the desired condition

---

**model** Union[str, Variant] required

Model to use for generating conditions

---

Returns:


**conditional** ConditionalGroup

Generated ConditionalGroup

---

Example:

---



```
# Generate conditional based on description - this will create conditions for
conditional= client.features.AutoConditional(
    "when the model talks about whales and penguins",
    Classes > Conditionals
    model=variant
)
# Get pirate feature
pirate_feature = client.features.search(
    "talk like a pirate", model=variant, top_k=1
)
# Make the model talk like a pirate when it talks about both whales and penguins
variant.set_when(conditional, {
    pirate_feature[0]: 0.9
})
```

## Best Practices

- Use conditional interventions to create context-aware steering behaviors
- Combine multiple conditions with logical operators for more precise control
- Handle aborted inferences gracefully in your application
- Test conditions thoroughly to ensure desired behavior
- Consider using AutoConditional for quick prototyping

## Classes

### ConditionalGroup

A group of conditions combined with logical operators.

Show Properties

### Conditional

A single conditional expression comparing features.

Show Properties

Context object containing information about the current inference state.

---

Classes > Conditionals

Show Properties

---

< Features

Auto Steer >

---

Powered by Mintlify



## Notebooks

# Decision Trees

Use features to label data and build decision trees

By inspecting feature activations on labelled datasets, we can build highly accurate and interpretable decision trees for classification tasks.

In this example, we classify financial news into positive and negative sentiments by building a decision tree.





## Notebooks

# Dynamic Prompts

By using Feature Inspection and Contrasting we can build a model that dynamically changes its instructions based on the user's prompt.

In this example, we'll build a model that dynamically changes its instructions based on if the user is asking for code to be written or not.







## Classes

# Features

Documentation for the Goodfire Features API

## Accessing Features

Before using the Features API, you'll need a model variant:

```
variant = Variant("meta-llama/Llama-3.3-70B-Instruct")
```

You can then access features through the client's features interface. For example, to search for features:

### Search for features

### Features

```
# Search for features related to a concept
features = client.features.search(
    "angry",
    model=variant,
    top_k=5
)

# Print the found features
print(features)
```

Or inspect feature activations in text:

### Inspect feature activations

### Features activated in text

```
# Analyze how features activate in text
inspector = client.features.inspect(
```



```
[
    {"role": "user", "content": "What do you think about pirates and ..."},
    {"role": "assistant", "content": "I think pirates are cool and whales a..."}
]
Classes > Features
model=variant

)

# Get top activated features
for activation in inspector.top(k=5):
    print(f"{activation.feature.label}: {activation.activation}")
```

The Features API provides methods for working with interpretable features of language models. Features represent learned patterns in model behavior that can be analyzed and modified.

## Methods

### neighbors()

Get the nearest neighbors of a feature or group of features.

#### Parameters:

**features** Feature | FeatureGroup **required**

Feature or group of features to find neighbors for

---

**model** str | VariantInterface **required**

Model identifier or variant interface

---

**top\_k** int default: 10

Number of neighbors to return

---

Returns: FeatureGroup

#### Example:

```
# Get neighbors of a feature
Classes > Features
neighbors = client.features.neighbors(
    feature,
    model="meta-llama/Llama-3.3-70B-Instruct",
    top_k=10
)

# Print neighbor labels
for neighbor in neighbors:
    print(neighbor.label)
```

## search()

Search for features based on semantic similarity to a query string.

### Parameters:

**query** str required

Search string to compare against feature labels

**model** str | VariantInterface required

Model identifier or variant interface

**top\_k** int default: 10

Number of features to return

**Returns:** `FeatureGroup` - Collection of matching features

### Example:

#### Search features

```
# Search for features related to writing style
features = client.features.search(
```



```
"formal writing style",  
model="meta-llama/Llama-3.3-70B-Instruct",  
top_k=10
```

```
)  
Classes > Features
```

```
# Print features  
for feature in features:  
    print(feature.label)
```

## inspect()

Analyzes how features are activated across the input messages.

### Parameters:

**messages** list[ChatMessage] required

Messages to analyze

**model** str | VariantInterface required

Model identifier or variant interface

**features** Feature | FeatureGroup | None

Optional specific features to analyze. If None, inspects all features.

**aggregate\_by** str default: "frequency"

Method to aggregate feature activations across tokens: - "frequency": Count of tokens where feature is active - "mean": Mean activation value across tokens - "max": Maximum activation value across tokens - "sum": Sum of activation values across tokens

**Returns:** `ContextInspector` - An inspector object that provides methods for analyzing and visualizing how features are activated in the given context.

### Example:

Inspect feature activations

```

 # Analyze how features activate in text
inspector = client.features.inspect(
    [
        {"role": "user", "content": "What do you think about pirates and whales"},
        {"role": "assistant", "content": "I think pirates are cool and whales are cute"}
    ],
    model=variant
)

# Get top activated features
for activation in inspector.top(k=5):
    print(f"{activation.feature.label}: {activation.activation}")

```

## contrast()

Identify features that differentiate between two conversation datasets.

Parameters:

**dataset\_1** list[list[ChatMessage]] required

First dataset of conversations

---

**dataset\_2** list[list[ChatMessage]] required

Second dataset of conversations

---

**model** str | VariantInterface required

Model identifier or variant interface

---

**top\_k** int default: 5

Number of top features to return for each dataset

---

**Returns:** `tuple[FeatureGroup, FeatureGroup]` - Two FeatureGroups containing:

Features steering towards dataset\_1



## Features steering towards dataset\_2

### Example:

Classes > Features

Get contrast features

```
# Compare formal vs informal conversations
dataset_1 = [[
    {"role": "user", "content": "Hi how are you?"},
    {"role": "assistant", "content": "I'm doing well..."}
]]

dataset_2 = [[
    {"role": "user", "content": "Hi how are you?"},
    {"role": "assistant", "content": "Arr my spirits be high..."}
]]

formal_features, informal_features = client.features.contrast(
    dataset_1=dataset_1,
    dataset_2=dataset_2,
    model="meta-llama/Llama-3.3-70B-Instruct",
    top_k=5
)
```

## rerank()

Rerank a set of features based on a query.

### Parameters:

**features** FeatureGroup **required**

Features to rerank

**query** str **required**

Query to rerank features by

**model** str | VariantInterface **required**

Model identifier or variant interface



**top\_k** int default: 10  
Classes > Features

Number of top features to return

Returns: `FeatureGroup`

Example:

#### Rerank

```
# Rerank features based on relevance to "writing style"
reranked = client.features.rerank(
    features=formal_features,
    query="writing style",
    model="meta-llama/Llama-3.3-70B-Instruct",
    top_k=10
)
```

## activations()

Retrieves feature activation values for each token in the input messages.

Parameters:

**messages** list[ChatMessage] required


Messages to analyze

**model** str | VariantInterface required

Model identifier or variant interface

**features** Feature | FeatureGroup | None

Optional specific features to analyze. If None, analyzes all features.

 **Returns:** `NDArray[np.float64]` - Sparse activation matrix of shape `[n_tokens, n_features]` where each element represents the activation strength of a feature at a specific token. Most values are zero due to sparsity.

[Classes](#) > [Features](#)

### Example:

#### Get activation matrix

```
# Get activation matrix for a conversation
matrix = client.features.activations(
    messages=[{"role": "user", "content": "Hello world"}],
    model="meta-llama/Llama-3.3-70B-Instruct"
)
```

## lookup()

Retrieves details for a list of features by their indices.

### Parameters:

**indices** `list[int]` **required**

List of feature indices to fetch

---

**model** `str | VariantInterface` **required**

Model identifier or variant interface

---

**Returns:** `dict[int, Feature]` - Mapping of feature index to Feature object

## list()

Retrieves details for a list of features by their UUIDs.

### Parameters:

**ids** `list[str]` **required**





Returns: `FeatureGroup` - Collection of Feature objects  
Classes → Features

## Classes

### Feature

A class representing a human-interpretable “feature” - a model’s conceptual neural unit. Features can be combined into groups and compared using standard operators.

Show Properties

### FeatureGroup

A collection of Feature instances with group operations.

Show Properties

### ConditionalGroup

Groups multiple conditions with logical operators.

Show Properties

### FeatureActivation

Represents the activation of a feature.

Show Properties

### ContextInspector

Analyzes feature activations in text.



Show Properties

Classes > **Features**  
< **Variants**

**Conditionals** >

---

Powered by Mintlify



## Notebooks

# Jailbreak Resistance

By using Feature Activations and Contrastive Search we can build a jailbreak resistant model.

Through this approach we were able to drastically lower the ability to jailbreak the model, using jailbreak prompts from the StrongREJECT dataset.





## Notebooks

# On Demand RAG

By using conditional interventions we can build a “on demand RAG” system.

When we see the user is asking about something that might need more data, we can abort the request and get more data from an external source.

In this example, we’ll build a model that can insert brand deal data into the response when the user is asking about products.

For example, if the user asks about drinks, and we sponsor Coca Cola, we can stop the request, get RAG data on brand deals and pass it back into the model.






## Get Started

# Quickstart

Get started with the Goodfire Ember SDK

 **Prerequisite:** You'll need a Goodfire API key to follow this guide. Get one through [our platform](#) or contact [support](#).

 Open in Colab

# Quickstart

Ember is a hosted API/SDK that lets you shape AI model behavior by directly controlling a model's internal units of computation, or **"features"**. With Ember, you can modify features to precisely control model outputs, or use them as building blocks for tasks like classification.

In this quickstart, you'll learn how to:

Find features that matter for your specific needs

Edit features to create model variants

Discover which features are active in your data

Save and load your model variants

### Code

```
pip install goodfire
```

You can get an API key through [our platform](#)



Get Started > Quickstart

## Initialize the SDK

### Code

```
import goodfire
client = goodfire.Client(api_key=GOODFIRE_API_KEY)
# Instantiate a model variant.
variant = goodfire.Variant("meta-llama/Llama-3.3-70B-Instruct")
```

Our sampling API is OpenAI compatible, making it easy to integrate.

### Generation Code

### Generation Output

```
for token in client.chat.completions.create(
    [{"role": "user", "content": "Hi, how are you?"}],
    model=variant,
    stream=True,
    max_completion_tokens=100,
):
    print(token.choices[0].delta.content, end="")
```

## Editing features to create model variants

### How to find relevant features for edits

There are three ways to find features you may want to modify:

**Auto Steer:** Simply describe what you want, and let the API automatically select and adjust feature weights

**Feature Search:** Find features using semantic search



Get Started > Quickstart

## Auto Steer

Auto steering automatically finds and adjusts feature weights to achieve your desired behavior. Simply provide a short prompt describing what you want, and autosteering will:

Find the relevant features

Set appropriate feature weights

Return a FeatureEdits object that you can set directly

### Auto Steer Code

### Auto Steer Output

```
edits = client.features.AutoSteer(  
    specification="be funny", # or your desired behavior  
    model=variant,  
)  
variant.set(edits)  
print(edits)
```

Now that we have a few funny edits, let's see how the model responds!

### Auto Steer Generation Code

### Auto Steer Generation Output

```
for token in client.chat.completions.create(  
    [{"role": "user", "content": "Tell me about pirates"}],  
    model=variant,  
    stream=True,  
    max_completion_tokens=120,  
):  
    print(token.choices[0].delta.content, end="")
```

The model automatically added puns/jokes, even though we didn't specify anything about comedy in our prompt.



Get Started > Quickstart

```
variant.reset()
```

Feature search helps you explore and discover what capabilities your model has. It can be useful when you want to browse through available features.

Feature Search Code

Feature Search Output

```
funny_features = client.features.search(
    "funny",
    model=variant,
    top_k=10
)
print(funny_features)
```

When setting feature weights manually, start with 0.5 to enhance a feature and -0.3 to ablate a feature. When setting multiple features, you may need to tune down the weights.

Feature set and generation code

Feature set and generation output

```
variant.set(funny_features[0], 0.6)
for token in client.chat.completions.create(
    [
        {"role": "user", "content": "tell me about foxes"}
    ],
    model=variant,
    stream=True,
    max_completion_tokens=100,
):
    print(token.choices[0].delta.content, end="")
```

Feel free to play around with the weights and features to see how the model responds.





Get Started > Quickstart

the group's centroid.

`neighbors()` helps you understand feature relationships beyond just their labels. It can reveal which features might work best for your intended model adjustments.

#### Nearest Neighbors Code

#### Nearest Neighbors Output

```
client.features.neighbors(  
    funny_features[0],  
    model=variant,  
    top_k=5  
)
```

Now, you can find more features that are similar to other features

#### Nearest Neighbors

#### Out

```
client.features.neighbors(  
    funny_features[2],  
    model=variant,  
    top_k=5  
)
```

## Contrastive Search

Contrastive search lets you discover relevant features in a data-driven way.

Provide two datasets of chat examples:

dataset\_1: Examples of behavior you want to avoid

dataset\_2: Examples of behavior you want to encourage

Examples are paired such that the first example in dataset\_1 contrasts the first example in dataset\_2, and so on.



This two-step process ensures you get features that are both:

Mechanistically useful (from contrastive search)

Aligned with your goals (from reranking)

Let's specify two conversation datasets. The first has a typical helpful assistant response and the second assistant replies in jokes.

#### Contrastive Search Code

#### Contrastive Search Output

```
variant.reset()
default_conversation = [
    [
        {
            "role": "user",
            "content": "Hello how are you?"
        },
        {
            "role": "assistant",
            "content": "I am a helpful assistant. How can I help you?"
        }
    ]
]
joke_conversation = [
    [
        {
            "role": "user",
            "content": "Hello how are you?"
        },
        {
            "role": "assistant",
            "content": "What do you call an alligator in a vest? An investigator"
        }
    ]
]
helpful_assistant_features, joke_features = client.features.contrast(
```



## Get Started > Quickstart

```
# Let's rerank to surface humor related features
joke_features = client.features.rerank(
    features=joke_features,
    query="funny",
    model=variant,
    top_k=5
)
joke_features
```

We now have a list of features to consider adding. Let's set some plausible-looking ones from `joke_features`.

Note that we could also explore removing some of the helpful\_assistant features.

### Contrastive Search Set Features Code

### Contrastive Search Set Features Output

```
variant.reset()
variant.set(joke_features[0,1], 0.6)
for token in client.chat.completions.create(
    [
        {"role": "user", "content": "Hello. Tell me about whales."}
    ],
    model=variant,
    stream=True,
    max_completion_tokens=100,
):
    print(token.choices[0].delta.content, end="")
```

## (Advanced) Conditional logic for feature edits

You can establish relationships between different features (or feature groups) using conditional interventions.

First, let's reset the variant and pick out the funny features.



Get Started > Quickstart

Now, let's find a features where the model is talking like a pirate.

#### Pirate Features Code

#### Pirate Features Output

```
pirate_features = client.features.search(
    "talk like a pirate",
    model=variant,
    top_k=3
)
print(pirate_features)
```

Now, let's set up behaviour so that when the model is talking like a pirate, it will be funny.

#### Conditional Logic Set Features Code

#### Conditional Logic Set Features Output

```
variant.set_when(pirate_features[1] > 0.75, {
    funny_features[0]: 0.7,
})
# The model will now try to be funny when talking about pirates
response = client.chat.completions.create(
    messages=[{"role": "user", "content": "talk like a pirate and tell me about"},
    model=variant
)
print(response.choices[0].message["content"])
```

Say we decide the model isn't very good at pirate jokes. Let's set up behavior to stop generation altogether if the pirate features are too strong.

#### Abort when Pirate Features are too strong

```
# Abort if pirate features are too strong
variant.abort_when(pirate_features > 0.75)
try:
```



Get Started > Quickstart

```
print("Generation aborted due to too much pirate content")
```

If you aren't sure of the features you want to condition on, use `AutoConditional` with a specified prompt to get back an automatically generated condition.

#### Auto Conditional Code

```
# Generate auto conditional based on a description. This will automatically
# choose the relevant features and conditional weight
conditional = client.features.AutoConditional(
    "pirates",
    model="meta-llama/Llama-3.3-70B-Instruct",
)

# Apply feature edits when condition is met
variant.set_when(conditional, {
    joke_features[0]: 0.5,
    joke_features[1]: 0.5
})
```

## Discover which features are active in your data

### Working with a conversation context

You can inspect what features are activating in a given conversation with the `inspect` API, which returns a `context` object.

Say you want to understand what model features are important when the model tells a joke. You can pass in the same joke conversation dataset to the `inspect` endpoint.

#### Inspect Code

#### Inspect Output



Get Started > Quickstart

```
)  
context
```

From the context object, you can access a lookup object which can be used to look at the set of feature labels in the context.

Lookup Code

Lookup Output

```
lookup = context.lookup()  
lookup
```

You can select the top  activating features in the context, ranked by activation strength. There are features related to jokes and tongue twisters, among other syntactical features.

Inspect Top Features Code

Inspect Top Features Output

```
top_features = context.top(k=10)  
top_features
```

You can also inspect individual tokens level feature activations. Let's see what features are active at the punchline token.

Token Activations Code

Token Activations Output

```
print(context.tokens[-3])  
token_acts = context.tokens[-3].inspect()  
token_acts
```

## (Advanced) Look at next token logits

Next Logits Code

Next Logits Output



Get Started > Quickstart

```
logits.logits
```

## Get feature activation vectors for machine learning tasks

To run a machine learning pipeline at the feature level (for instance, for humor detection) you can directly export features using `client.features.activations` to get a matrix or retrieve a sparse vector for a specific `FeatureGroup`.

### Matrix Code

### Matrix Output

```
activations = client.features.activations(  
    messages=joke_conversation[0],  
    model=variant,  
)  
activations
```

### Vector Code

### Vector Output

```
top_features.vector()
```

## Inspecting specific features

There may be specific features whose activation patterns you're interested in exploring. In this case, you can specify features such as *humor\_features* and pass that into the `features` argument of `inspect`.

### Humor Features Code

### Humor Features Output

```
humor_features = client.features.search("jokes and humor", model=variant, top_k=10)  
humor_features
```



## Get Started > Quickstart

```
context = client.features.inspect(  
    messages=joke_conversation[0],  
    model=variant,  
    features=humor_features  
)  
context
```

Now you can retrieve the top k activating *humor features* in the `context`. This might be a more interesting set of features for downstream tasks.

### Top Features Code

### Top Features Output

```
humor_feature_acts = context.top(k=5)  
humor_feature_acts
```

## Save and load your model variants

You can serialize a variant to JSON format for saving.

### Save Variant Code

### Save Variant Output

```
variant.reset()  
variant.set(pirate_features[1], 0.9)  
variant_json = variant.json()  
variant_json
```

And load a variant from JSON format.

### Load Variant Code

### Load Variant Output

```
loaded_variant = goodfire.Variant.from_json(variant_json)  
loaded_variant
```





Get Started > Quickstart

```
for token in client.chat.completions.create(
    [
        {"role": "user", "content": "tell me about whales"}
    ],
    model=loaded_variant,
    stream=True,
    max_completion_tokens=150,
):
    print(token.choices[0].delta.content, end="")
```

## Using OpenAI SDK

You can also work directly with the OpenAI SDK for inference since our endpoint is fully compatible.

### Install OpenAI Code

```
!pip install openai --quiet
```

### OpenAI SDK Code

```
from openai import OpenAI

oai_client = OpenAI(
    api_key=GOODFIRE_API_KEY,
    base_url="https://api.goodfire.ai/api/inference/v1",
)

response = oai_client.chat.completions.create(
    messages=[
        {"role": "user", "content": "who is this"},
    ],
    model=variant.base_model,
    extra_body={"controller": variant.controller.json()},
```



---

Get Started > Quickstart

! For more advanced usage and detailed API reference, check out our [SDK reference](#) and [example notebooks](#).

< Introduction

Decision Trees >

---

Powered by Mintlify



## Notebooks

# Removing Knowledge

By creating a new Variant with certain features removed, we can remove knowledge from the model.

In this example, we'll create a variant which doesn't know about famous celebrities.





## Notebooks

# Sorting by Features

By using Feature Activations we can sort data by the features that are most relevant to a given query.

In this example, we'll sort Elon Musk's tweets by the sarcasm feature.



< Removing Knowledge



## Classes

# Variants

Documentation for working with model variants

Variants are edits to a model that allow you to modify model behavior by adjusting feature activations and defining conditional behaviors.

## Creating Variants

### Basic Usage

Create a variant by instantiating the `Variant` class with a base model:

Code

Output

```
from goodfire import Variant

# Create a variant from a base model
variant = Variant("meta-llama/Llama-3.3-70B-Instruct")
print(variant)
```

## Adding features to a variant

Adding features to a variant

Edited variant

```
from goodfire import Variant

# Create a variant from a base model
variant = Variant("meta-llama/Llama-3.1-8B-Instruct")

# Search for a feature to modify
```



```
feature = client.features.search("formal writing style", model=variant, top_k=1)

# Set feature modifications
variant.set(feature, 0.5) # Value typically between -1 and 1
```

Classes > Variants

## Conditional Controls

You can create variants that respond dynamically to feature activations:

### Conditional Controls

```
# First get some features to work with
whale_feature = client.features.search("whales", model=variant, top_k=1)[0]
pirate_feature = client.features.search("pirate speech", model=variant, top_k=1)[0]

# Activate pirate features when whale features are detected
variant.set_when(whale_feature > 0.75, {
    pirate_feature: 0.5
})

# Abort generation if certain features are too strong
variant.abort_when(whale_feature > 0.9)

# Custom handler when condition is met
def my_handler(context):
    print(f"Whale feature activated with strength: {context.activations[whale_feature]})
```

## Methods

### set()

Set feature modifications. This method is overloaded to handle different input types.

Signatures:

```
def set(self, feature: Feature | FeatureGroup, value: float)
```



```
def set(self, edits: dict[Feature, float] | FeatureEdits)
```

## Parameters > Variants

**feature** Feature | FeatureGroup

Single feature or feature group to modify. Required when using the first signature.

**value** float

Modification value (typically between -1 and 1). Required when using the first signature.

**edits** dict[Feature, float] | FeatureEdits

Dictionary of features and their values, or a FeatureEdits object. Required when using the second signature.

## Examples:

```
# Single feature modification
formal_feature = client.features.search("formal writing style", model=variant,
variant.set(formal_feature, 0.5)

# Multiple features at once using dict
casual_feature = client.features.search("casual writing style", model=variant,
variant.set({
    formal_feature: 0.5,
    casual_feature: -0.3
}))

# Using FeatureEdits object
edits = client.features.AutoSteer("be funny", model=variant)
variant.set(edits)
```

## set\_when()

Define conditional feature modifications.

## Parameters:



**condition** ConditionalGroup required

Classes > Variants

Condition that triggers the modifications

---

**values** dict[Feature, float] | FeatureEdits required

Feature modifications to apply when condition is met

---

### Example:

#### set\_when example

---

```
whale_feature = client.features.search("whales", model=variant, top_k=1)[0]
pirate_feature = client.features.search("pirate speech", model=variant, top_k=1)[0]
# Activate pirate features when whale features are detected
variant.set_when(whale_feature > 0.3, {
    pirate_feature: 0.5
})
```

## abort\_when()

Abort generation when a condition is met.

## Parameters:

**condition** ConditionalGroup required

Condition that triggers the abort

---

### Example:

#### Abort example

---

```
inappropriate_feature = client.features.search("inappropriate content", model=variant, top_k=1)[0]
# Abort if inappropriate content is detected
variant.abort_when(inappropriate_feature > 0.8)
```





```
try:
    response = client.chat.completions.create(
        messages=[{"role": "user", "content": "..."}],
        model=variant,
        stream=False
    )
except goodfire.exceptions.InferenceAbortedException:
    print("Aborted because of inappropriate content")

print(variant)
```

## reset()

Remove all feature modifications.

Example:

Reset variant

```
variant.reset()
```

## clear()

Remove modifications for specific features.

Parameters:

**feature** Feature | FeatureGroup **required**

Feature(s) to clear modifications for

Example:

Clear particular feature

```
variant.clear(feature)
```

# Serialization

Variants can be serialized to and from JSON:

Classes > Variants

## Storing and loading variants

```
# Save variant to JSON
variant_json = variant.json()

# Load variant from JSON
loaded_variant = Variant.from_json(variant_json)
```

## Using with OpenAI SDK

Variants are compatible with the OpenAI SDK:

### Using with OpenAI SDK

```
from openai import OpenAI

client = OpenAI(
    api_key="YOUR_GOODFIRE_API_KEY",
    base_url="https://api.goodfire.ai/api/inference/v1"
)

response = client.chat.completions.create(
    messages=[{"role": "user", "content": "Hello"}],
    model=variant.base_model,
    extra_body={"controller": variant.controller.json()}
)
```

## Classes

### VariantMetaData

Metadata about a model variant.

Properties:

 `name` str

Name of the variant

[Classes](#) > [Variants](#)

`base_model` str

Base model identifier

`id` str

Unique identifier for the variant

[Chat Completions](#)

[Features](#)

Powered by Mintlify